

Lecture 19: Dimensionality Reduction

Dimensionality reduction

- Representing data using **a smaller number of variables** while preserving the "interesting" structure of the data
- Purposes:
 - Visualisation
 - Computational efficiency in a pipeline
 - Data compression or statistical efficiency in a pipeline
- Results in loss of information in general
 - Trick: ensure that most of the "interesting" information (signal) is preserved (while what is lost is mostly noise)

Principal component analysis (PCA)

- Popular method for dimensionality reduction and data analysis
- Aim: find **a new coordinate system** s.t. most of the **variance is concentrated** along the first coordinate, then most of the remaining variance along the second coordinate, etc.
- Dimensionality reduction is based on **discarding coordinates** except the first $l < m$

Formulating the problem

- Projection of u on v : $u_v = u \cdot v / \|v\|$
 - If $\|v\| = 1$, $u_v = u \cdot v$
- Vector v can be considered as a candidate **coordinate axis**
 - and u_v the coordinate of point u (on the new coordinate axis)
- Data transformation
 - Projecting all data points to a new coordinate axis p_1
 - result: $X'p_i$
 - Where $\|p_i\| = 1$
 - X has original data points in columns
- Sample covariance matrix:
 - For centered (mean subtracted) matrix X
 - $\Sigma_X = \frac{1}{n-1} X'X$

PCA

- Objective:
 - assume the data is centered
 - find p_1 to maximise variance along this PC: $p_1' \Sigma_X p_1$ subject to $\|p_1\| = 1$
- Constrained problem -> Use Lagrange multipliers
- Solution: p_1 is the eigenvector with corresponding λ_1 being the **max** eigenvalue (of covariance matrix Σ_X)
 - Variance captured by PC1: $\lambda_1 = p_1' \Sigma_X p_1$
- (Spectrum of a matrix is a set of its eigenvalues)
- Choose dimensions to keep from "knee" in scree plot


Efficient algorithm for PCA

- Setting p_i as all eigenvectors of the **centered** data covariance matrix Σ_X in decreasing eigenvalue order
- Lemma: a real symmetric $m \times m$ matrix has m real eigenvalues and corresponding eigenvectors are orthogonal
- Lemma: a PSD matrix further has non-negative eigenvalues

Linear regression v.s. PCA

- Another view of PCA: s — *dim* plane minimising residual sum squares to data
- It turns out:
 - PCA chooses the **direction** to be a hyperplane that minimise these errors (RSS)
 - Since variance and squared distance have something in common (both sum of squares)

PCA vs. Linear regression



Additional effect of PCA

- Consider candidate axes i and $(i + 1)$, if there is correlation between them
 - This means that axis i can be rotated further to capture more variance
- PCA should end up finding new axes (transformation) s.t. the transformed data is uncorrelated

Non-linear data and kernel PCA

- Low dimensional approximation need not be linear
- Kernel PCA: **map** data to feature space, **then** run **PCA**
 - Express PC in terms of data points
 - Solution uses $X'X$ that can be kernelised:
 - $(X'X)_{ij} = K(x_i, x_j)$
 - Solution strategy differs from regular PCA
 - Modular: Changing kernel leads to a different feature space transformation